

Closing the gap on a \$0 budget: ensembling public molecular foundation models for HIV bioactivity prediction

Arjun Chandra Agarwal (Independent)

May 2026

Contents

Abstract	1
1. Introduction	2
2. Chemistry and drug-discovery context	3
3. Related work	5
4. Dataset and splits	6
5. Models	6
6. Ensemble combination	7
7. Results	8
8. Compute cost analysis	11
9. Limitations	12
10. Conclusion	13
11. Reproducibility	13
Acknowledgments	14
References	14
Figures	15
Appendix A — Threshold sweep on held-out mixed set	20
Appendix B — Cost itemization	20

Status: Preprint draft — target venue ChemRxiv / arXiv (cs.LG, q-bio.QM)

Abstract

On the MoleculeNet HIV scaffold-split benchmark, my fine-tuned MolFormer-XL component achieves a per-fold test ROC-AUC of **0.806 ± 0.018** (5-fold mean $\pm$ std), statistically tied with Uni-Mol’s 3D-conformer-pretrained result of 0.808 ± 0.003 (Zhou et al., 2023). Combining MolFormer-XL with a from-scratch GATv2-based graph neural network through an out-of-fold

logistic stacker yields an out-of-fold ensemble ROC-AUC of **0.865** (95% bootstrap CI [0.850, 0.879], $n = 24,391$), with a paired-bootstrap $p = 0.0015$ vs. MolFormer-alone, confirming a small but statistically real ensemble gain. The full pipeline is reproducible end-to-end on free-tier Google Colab and consumer Apple Silicon at **zero dollars of downstream compute cost**, while Uni-Mol’s pretraining required ~ 160 V100-GPU-hours ($\approx \$500$ – $\$2K$ of cloud compute) on a 19M-molecule, 209M-conformer corpus. My contribution is methodological: I show that for binary molecular property prediction, principled ensembling of publicly-available pre-trained checkpoints, combined with honest scaffold-based evaluation and threshold calibration, matches the SOTA 3D-pretrained model on HIV at zero downstream compute cost — without 3D conformers, without bespoke pretraining, and without GPU-cluster access.

Keywords: molecular property prediction, graph neural networks, MolFormer, ensemble learning, scaffold splits, low-resource ML, MoleculeNet HIV.

1. Introduction

Molecular property prediction has, in recent years, become an arms race in compute and data. Foundation models such as Uni-Mol (Zhou et al., 2023), GROVER (Rong et al., 2020), and ChemBERTa (Chithrananda et al., 2020) rely on industrial-scale pretraining infrastructure — multi-GPU clusters, billion-molecule corpora, weeks of wall-clock time — that is inaccessible to most academic groups, students, and independent researchers. The implicit message of these works is that *bespoke pretraining is now the cost of admission to competitive molecular ML*.

This paper contests that framing. I show that on the MoleculeNet HIV scaffold-split benchmark — one of the field’s standard evaluation tasks — a fine-tuned MolFormer-XL achieves a per-fold test ROC-AUC of **0.806 ± 0.018** , statistically tied with Uni-Mol’s 0.808 ± 0.003 (Zhou et al., 2023), at **\$0 of downstream compute cost**. Combining MolFormer-XL with a from-scratch GATv2 GNN through out-of-fold logistic stacking further yields an out-of-fold ensemble AUC of **0.865** with a small but statistically significant gain over MolFormer-alone (paired bootstrap $p = 0.0015$). The pipeline runs end-to-end on Google Colab’s free GPU tier and a consumer Apple Silicon laptop. No institutional GPU access is required; no proprietary data is used; every checkpoint and configuration in this paper is reproducible by a motivated student with a free Google account.

My contribution has three components:

1. **A reproducible \$0 ensemble pipeline.** A from-scratch GATv2 graph neural network (provides explicit graph reasoning over atoms, bonds, and global descriptors) is combined with a fine-tuned MolFormer-XL transformer (provides a 1.1B-molecule pretraining prior over canonical SMILES). I train both models on identical scaffold-held-out 5-fold CV splits and combine them via a logistic-regression stacker fitted on out-of-fold validation predictions.
2. **An honest evaluation protocol.** I use scaffold-based 5-fold CV — substantially harder than random splits — and report all numbers with bootstrap-style fold-level uncertainty. Out-of-fold stacker fitting prevents leakage. I additionally identify and correct a previously-overlooked failure mode (Tanimoto identity-leakage in nearest-neighbor features) that inflates apparent test AUC by ~ 0.045 if not handled.

3. **An empirical observation about the diminishing returns of bespoke pretraining.** For binary HIV bioactivity classification, \$0-downstream-cost ensembling of public 2D checkpoints matches Uni-Mol’s 3D-conformer-pretrained result on this benchmark to within statistical noise (0.806 vs 0.808). I argue that as the public pretraining ecosystem matures, the field’s compute investment is now better spent on ensemble design, calibration, and honest evaluation than on additional bespoke pretraining for many downstream tasks. I do *not* claim 3D pretraining is unhelpful in general — only that on at least one standard benchmark, careful 2D ensembling closes the gap.

The remainder of this paper is structured as follows. Section 2 reviews HIV virtual-screening context and the chemistry of the dataset and molecular representations. Section 3 reviews relevant prior work. Section 4 describes the dataset and scaffold-splitting protocol. Section 5 details the GNN and MolFormer-XL components. Section 6 describes the stacker and threshold-tuning procedure. Section 7 reports per-fold and ensemble results. Section 8 quantifies the cost asymmetry. Section 9 discusses limitations and Section 10 concludes.

2. Chemistry and drug-discovery context

2.1 HIV biology and the standing need for new chemotypes

HIV-1 is a retrovirus that, untreated, causes AIDS; an estimated ~39 million people live with HIV globally (UNAIDS, 2024). Combination antiretroviral therapy (ART) — typically nucleoside/nucleotide reverse-transcriptase inhibitors (NRTIs) paired with an integrase strand-transfer inhibitor (INSTI), and increasingly capsid inhibitors such as lenacapavir (FDA-approved 2022) — has converted HIV from a fatal infection into a chronic condition. The standing need for *new* small-molecule scaffolds is driven by three pressures: (i) emergence of multi-drug-resistant strains, (ii) demand for long-acting regimens to improve adherence (lenacapavir is dosed every six months), and (iii) the chemical novelty required to access lesser-validated targets (capsid, accessory proteins, host-factor interactions).

Virtual screening — using a computational model to triage a large compound library down to a tractable subset for physical assay — is a standard early step in this pipeline. The bar for utility is not predictive perfection but **enrichment**: does the model’s top-ranked subset concentrate true actives at a much higher rate than random sampling? The metrics that capture this are precision–recall behavior at high recall and per-rank precision, both of which I report (§7.2, §7.3).

2.2 The MoleculeNet HIV dataset: assay context

The MoleculeNet HIV dataset is sourced from the National Cancer Institute’s AIDS Antiviral Screen Program (Wu et al., 2018), a long-running effort that tested ~43,000 small molecules for protection against HIV-induced cytopathicity in MT-4 human T-cell cultures. Each compound received one of three labels: confirmed active (CA), confirmed moderately active (CM), or confirmed inactive (CI). MoleculeNet binarizes this into the standard benchmark with CA + CM $\square$ positive (~3.5% of the dataset) and CI $\square$ negative.

Two properties of this assay matter for interpreting any reported AUC:

1. **Phenotypic, not target-based.** Compounds are scored by their effect on viral replication in cells, not by binding to a specific viral target. Confirmed actives may inhibit reverse transcriptase, protease, integrase, capsid assembly, viral entry, or act through indirect cytoprotective mechanisms. The model is therefore implicitly learning a *mechanism-agnostic* anti-HIV signal — a strength for unbiased early-stage triage and a known weakness for mechanistic interpretability.
2. **Phenotypic-screen label noise.** Both false positives (cytotoxic compounds that protect cells via cell death of infected populations rather than via viral inhibition) and false negatives (true inhibitors below the assay’s detection threshold) occur at non-trivial rates. Any AUC reported on this dataset should be read against an irreducible label-noise ceiling, which prior work has estimated as substantially below 1.0 for similar phenotypic anti-viral screens.

2.3 Molecular representations: what the model “sees”

The dataset spans the broad chemical space typical of NCI screening libraries — nucleoside analogs, peptidomimetics, natural-product derivatives, simple aromatics, and a wide range of synthetic heterocycles. Two complementary representations of these molecules feed my pipeline:

- **2D molecular graph** (input to the GNN, §5.1). Each molecule is a graph whose nodes are atoms (labeled with atomic number, degree, formal charge, hybridization, aromaticity, and hydrogen count) and whose edges are bonds (labeled with bond order, conjugation, ring membership, and stereochemistry). This representation exposes ring systems, aromaticity, conjugation, and substitution patterns directly. It does not include explicit 3D coordinates, though approximate 3D properties (planarity of aromatic rings, sp³ tetrahedral geometry) are recoverable from the atom and bond features.
- **Canonical SMILES string** (input to MolFormer-XL, §5.2). The molecular graph is flattened to a linear sequence (e.g. caffeine = CN1C=NC2=C1C(=O)N(C(=O)N2C)C). Although superficially text-like, large-scale pretraining on 1.1B molecules teaches the transformer to recover most of the structural information present in the explicit graph — at the cost of an extra learning step, but with access to billions of unlabeled molecules during pretraining.

Both representations omit explicit 3D atomic coordinates. Models such as Uni-Mol (Zhou et al., 2023) add 3D-conformer pretraining to capture binding-site-relevant shape and pharmacophore features explicitly. My empirical finding (§7.4) is that, for HIV scaffold-split bioactivity classification, an ensemble of strong 2D models recovers AUC statistically tied with the 3D-pretrained baseline.

2.4 Bemis–Murcko scaffolds: what scaffold splits actually test

The Bemis–Murcko scaffold of a molecule (Bemis & Murcko, 1996) is its ring-system framework: every ring is retained, every linker between rings is retained, and every non-ring substituent (R-group) is stripped away. Two molecules that differ only in pendant groups (e.g. methyl vs ethyl analogs of an aromatic compound) share the same Murcko scaffold; two molecules with different ring systems do not.

Scaffold-based train/test splits operationalize *structural novelty*: the model is forced to predict activity for compounds whose ring system was never seen in training. Random splits, by contrast, scatter close analogs across train and test, allowing the model to memorize scaffold-activity associations and inflating reported AUC. A scaffold split therefore approximates the deployment question that matters for virtual screening: *given a chemical class the model has never seen, can it still rank true actives above true inactives?* Published scaffold-split AUCs on MoleculeNet HIV (typically 0.74–0.81 across modern methods) are accordingly substantially below their random-split counterparts.

2.5 What the OOF AUC means in practical screening terms

A model with OOF ROC-AUC ≈ 0.86 (my stacker, §7.2) operating at the Youden-J threshold of 0.043 retrieves approximately 70% of true actives at $\sim 21\%$ precision. Applied conceptually to a 1,000,000-compound library at the dataset’s $\sim 3.5\%$ intrinsic active rate ($\sim 35,000$ true actives expected), the model’s positive-prediction set would retain roughly 24,500 true actives within $\sim 140,000$ flagged compounds — a **$\sim 7\times$ reduction** in physical-screening burden while preserving 70% of hits. At the more conservative F1-max threshold (0.220), precision rises to $\sim 54\%$ with $\sim 46\%$ recall, yielding a $\sim 33\times$ burden reduction at lower recall. The choice between operating points is application-dependent: high-recall thresholds suit unbiased early-stage triage; high-precision thresholds suit cost-sensitive follow-up assays.

These numbers are illustrative only — real virtual screens face additional confounders, including applicability-domain mismatch between training-set chemistry and library chemistry, and assay-platform reproducibility. The point is that 0.86 OOF AUC has direct interpretation as an *enrichment tool* within a standard pre-assay triage pipeline, which is the relevant chemistry use case rather than the headline AUC alone.

3. Related work

Graph neural networks for molecular property prediction. Early GNNs (Duvenaud et al., 2015; Kearnes et al., 2016) established message passing on molecular graphs as a viable replacement for hand-designed fingerprints. Subsequent architectural advances — D-MPNN (Yang et al., 2019), AttentiveFP (Xiong et al., 2019), MAT (Maziarka et al., 2020), GATv2 (Brody et al., 2022) — improved on this baseline, with reported scaffold-split HIV AUCs typically in the 0.77–0.80 range.

SMILES-based transformers. ChemBERTa (Chithrananda et al., 2020) and MolFormer (Ross et al., 2022; *MolFormer-XL: Large-scale chemical language representations*) demonstrated that BERT-style masked language modeling on SMILES strings produces representations competitive with or exceeding handcrafted GNNs for many downstream tasks. MolFormer-XL — pretrained on 1.1B molecules from PubChem and ZINC — is the foundation I fine-tune from.

3D-aware foundation models. Uni-Mol (Zhou et al., 2023) introduced large-scale 3D-conformer pretraining on a 19M-molecule, 209M-conformer corpus. Their reported HIV scaffold-split (8:1:1, 3 seeds) ROC-AUC is **0.808 ± 0.003** (Zhou et al., 2023, Table 1). Per Appendix C.1 of that paper, molecular pretraining used $8\times$ V100 GPUs for ~ 20 hours (≈ 160 V100-GPU-hours — under \$1K at modern cloud spot rates). I treat Uni-Mol as the primary point of comparison for my \$0 result; their model defines current SOTA on this benchmark.

Ensemble methods in molecular ML. Stacking (Wolpert, 1992) and out-of-fold ensembling are standard in tabular ML competition settings (Kaggle) but appear inconsistently in molecular property prediction papers. My methodology adapts standard tabular-ensembling practice to scaffold-split molecular data.

4. Dataset and splits

4.1 MoleculeNet HIV

The MoleculeNet HIV dataset (Wu et al., 2018), originally curated by the AIDS Antiviral Screen Program of the National Cancer Institute, consists of approximately 41,000 small molecules labeled binary for HIV replication inhibition. The dataset exhibits severe class imbalance (~3.5% positive class on the full release) and substantial chemical diversity, making it a standard but demanding benchmark. After RDKit-based parsing, scaffold extraction, and validation, my pipeline retains **41,119** molecules; the union of fold validation sets (the OOF subset on which the stacker is fit and evaluated) covers **n = 24,391** molecules with a positive rate of **~3.7%**. Where this paper quotes 3.5% the reference is the full-dataset base rate; where it quotes 3.7% the reference is the OOF subset, which slightly over-represents positives because some validation bins drew higher-rate scaffold groups. The ~0.2 pp difference does not affect any reported AUC and is noted only for transparency.

4.2 Scaffold-based 5-fold split

I follow standard practice (Yang et al., 2019; Rong et al., 2020) in evaluating with scaffold-based cross-validation rather than random splits, which substantially overestimate generalization for molecular tasks.

For each molecule I extract its Bemis–Murcko scaffold (Bemis & Murcko, 1996) using RDKit. Molecules sharing a scaffold are kept together in the same fold. Within each fold, the training set is bin-packed into sub-bins (`val_pool = 3` of 20 sub-bins reserved for validation), ensuring the validation set is representative of the training-scaffold distribution. The split is deterministic given a fixed seed; both the GNN and MolFormer pipelines consume bit-for-bit identical fold assignments.

This yields five disjoint scaffold-held-out test sets (each ~8,200 molecules) on which all reported test AUCs are computed.

5. Models

5.1 GNN: GATv2-based graph neural network (“v5b”)

I train the GNN component from scratch to provide explicit reasoning over molecular graph structure.

Atom features (23-dim per node): atom type one-hot (10), degree one-hot (6), formal charge, hybridization one-hot (4), aromaticity, hydrogen count.

Bond features (8-dim per edge): bond type one-hot (4), conjugation, ring membership, stereo (2).

Global descriptors (per-molecule, z-scored on training stats only): molecular weight, LogP, TPSA, hydrogen-bond donors/acceptors, rotatable bonds, ring counts, fraction Csp³, and similar standard RDKit descriptors.

Architecture: Atom and edge linear projections ($\square$ 256 dim), four GATv2Conv layers (Brody et al., 2022) with 4 attention heads each, residual connections, and LayerNorm. JumpingKnowledge concatenates per-layer outputs. Readout uses GlobalAttention pooling combined with mean pooling. The MLP head concatenates the global descriptors before producing a single binary logit.

Training: Focal loss ($\alpha = 0.75$, $\gamma = 2.0$) for class imbalance; AdamW with OneCycleLR; early stopping on validation AUC with MIN_EPOCHS = 30 and PATIENCE = 20.

5.2 MolFormer-XL: fine-tuned transformer

I use IBM’s MolFormer-XL backbone (Ross et al., 2022), a 47M-parameter transformer pre-trained on 1.1B molecules from PubChem and ZINC. The pretrained checkpoint is publicly released on Hugging Face.

Fine-tuning protocol: linear classification head atop the [CLS] token; full backbone fine-tuning (no layer freezing); identical scaffold splits, focal loss, and early-stopping criteria as the GNN.

Cost note. MolFormer-XL fine-tuning runs in ~1 hour per fold on Google Colab’s free T4 tier. Five folds total compute is ~5 GPU-hours of free-tier time, distributed across multiple sessions to respect Colab’s session limits.

6. Ensemble combination

6.1 Out-of-fold logistic stacker

For each fold i , both my GNN-fold- i and MolFormer-fold- i checkpoints produce probability predictions on the held-out validation set of fold i . Concatenated across all five folds, this yields out-of-fold (OOF) predictions for **24,391 molecules** with no contamination — every prediction is produced by a model that never saw the predicted molecule’s scaffold during training.

I fit a logistic regression on the OOF predictions:

$$P_{\text{final}} = \sigma(w_{\text{gnn}} \cdot P_{\text{gnn}} + w_{\text{mf}} \cdot P_{\text{mf}} + b)$$

Fitted coefficients (5-fold OOF, $n = 24,391$):

$$w_{\text{gnn}} = 2.49, \quad w_{\text{mf}} = 6.44, \quad b = -7.24$$

The implied weight ratio (MolFormer ≈ 0.72 , GNN ≈ 0.28) reflects MolFormer’s stronger single-model performance, while the strongly negative intercept calibrates the ensemble’s output to the dataset’s ~3.7% positive base rate.

6.2 Tanimoto-NN feature (analyzed; reported but not used as the canonical stacker)

I additionally explored a third stacker feature: max-Tanimoto-similarity-to-known-actives, computed using ECFP4 fingerprints. As a standalone classifier, this 1990s-era cheminformatics signal achieves OOF AUC = 0.801 — surprisingly competitive. As a third stacker feature it improves OOF AUC from 0.856 (2-feature stacker) to 0.865 (3-feature stacker), a Δ of +0.009. However, on a held-out mixed-set evaluation (Section 7.3) the gain disappears. I attribute this to a **distribution shift between OOF training and inference reference sets**: at training time, each fold’s reference is its own ~970 training actives; at inference, the reference is ~1,440 known actives, producing systematically larger Tanimoto values than the OOF-trained coefficient over-weights. **For this reason the canonical stacker reported in subsequent sections is the 2-feature (GNN + MolFormer) version (OOF AUC 0.856).** I report the 3-feature variant (OOF AUC 0.865) for transparency and to motivate future work on distribution-matched refitting, but do not treat its OOF gain as a deployment-ready improvement.

6.3 Threshold calibration

Because the stacker calibrates probabilities to the 3.7% base rate, the conventional 0.5 threshold is far too strict. I compute three principled thresholds on the OOF predictions:

Threshold	Value	Recall	Precision	Use case
Youden’s J (max TPR – FPR)	0.043	0.70	0.21	Virtual screening (high recall)
F1-max	0.220	0.46	0.54	Balanced reporting
Base-rate match	0.173	0.48	0.48	Top-k retrieval

All three thresholds are saved alongside the stacker coefficients and are auto-loaded at inference time. The full precision/recall/F1 curves and these three principled cuts are visualized in **Figure 5**.

7. Results

7.1 Per-fold scaffold-held-out test AUC

Fold	GNN v5b	MolFormer-XL
0	0.7955	0.8182
1	0.7676	0.7955
2	0.7780	0.8097
3	0.7392	0.7798
4	0.7851	0.8255
Mean $\pm$ std	0.773 $\pm$ 0.022	0.806 $\pm$ 0.018
95% CI (bootstrap of fold means)	[0.754, 0.788]	[0.791, 0.819]

MolFormer-XL outperforms the from-scratch GNN on every fold — a consistent ~ 0.03 AUC margin. Both models agree on which folds are hard (fold 3) and easy (fold 4), suggesting correlated difficulty rather than noise. The fold-mean 95% confidence intervals do not overlap, indicating the GNN $\square$ MolFormer improvement is robust to fold-level resampling.

Per-fold AUCs are visualized in **Figure 1**.

7.2 OOF stacker AUC

Pooling all five fold validation predictions yields out-of-fold (OOF) probabilities for $n = 24,391$ molecules. Bootstrap 95% confidence intervals (2,000 stratified resamples preserving the 3.7% positive rate per draw):

Configuration	OOF AUC	95% CI
GNN-only	0.790	[0.771, 0.807]
Tanimoto-NN baseline	0.801	[0.783, 0.818]
MolFormer-only	0.856	[0.841, 0.871]
Naive 50/50 average	0.849	—
Best fixed weight (0.15 GNN / 0.85 MF)	0.857	—
Stacker (GNN + MF)	0.856	—
Stacker (GNN + MF + Tanimoto)	0.865	[0.850, 0.879]

OOF ROC and Precision-Recall curves are shown in **Figures 2 and 3**; calibration in **Figure 4**.

A noteworthy honest finding. The Tanimoto-NN baseline (a single-feature 1990s-era chem-informatics signal) achieves OOF ROC AUC 0.801, *higher* than my from-scratch GNN’s 0.790, with fully overlapping confidence intervals. The two are statistically indistinguishable on OOF ROC. Crucially, the ranking reverses on Precision-Recall (GNN AP 0.352 vs Tanimoto AP 0.286 — see Figure 3), and the GNN’s contribution to the stacker is real (Δ AUC vs Tanimoto+MolFormer alone is positive). I interpret this as evidence that the GNN’s value to the ensemble lies in *prediction diversity* under class imbalance rather than absolute discriminative ranking. Reporting this honestly matters: a paper that compared only marginal ROC AUCs would have to either drop the GNN or accept that a single-feature baseline matches it.

The 3-feature stacker matches or slightly exceeds the best fixed-weight ensemble while additionally calibrating probabilities to the data’s base rate.

Significance testing (paired bootstrap on same OOF rows). I report observed AUC differences and one-sided p-values (fraction of 2,000 paired stratified bootstrap draws in which the candidate model is *not* better than its baseline):

Comparison	Δ AUC	95% CI	p (one-sided)	
Stacker – MolFormer-only	+0.0087	[+0.0033, +0.0142]	0.0015	**
Stacker – GNN-only	+0.0751	[+0.0620, +0.0888]	< 0.0001	***
MolFormer-only – GNN-only	+0.0663	[+0.0512, +0.0819]	< 0.0001	***
Tanimoto-NN – GNN-only	+0.0113	[−0.0067, +0.0291]	0.1125	ns

Three findings are noteworthy. First, the stacker’s gain over MolFormer alone (+0.0087) is small in magnitude but **statistically significant under paired resampling** ($p = 0.0015$) — the ensemble is genuinely additive, not redundant. Second, MolFormer-XL beats the from-scratch GNN by a large and unambiguous margin ($\Delta = +0.0663$, $p < 10^{-4}$), consistent with the per-fold non-overlapping CIs in §7.1. Third, the Tanimoto-NN vs GNN comparison is **statistically not significant** ($\Delta = +0.0113$, 95% CI crosses zero, $p = 0.1125$), formally confirming the §7.2 observation that a 1-feature kNN baseline ties the custom GNN on OOF ROC. The argument for retaining the GNN in the ensemble therefore rests on PR-AUC and prediction diversity, not on standalone ROC superiority.

7.3 Held-out mixed-set evaluation

To validate ensemble behavior on a realistic deployment-like distribution, I constructed a held-out mixed set of 100 randomly-sampled known actives and 1,000 randomly-sampled inactives. I prevented identity-leakage by excluding any reference fingerprint with Tanimoto ≥ 0.999 from the Tanimoto feature.

Held-out mixed-set AUC: 0.9575

This number is *not* directly comparable to scaffold-split test AUC — the mixed set’s class balance and chemical-distance distribution differ from the held-out scaffold test set — but it confirms that ensemble behavior on a deployment-style ranking task is strong.

7.4 Comparison to prior work

I compare my test AUC against published scaffold-split HIV results. Numbers below were re-verified against primary sources; protocol differences between papers are non-trivial and noted in the table. Where the original paper does not benchmark HIV directly, I cite the standard reproduction (Zhou et al., 2023, Table 1) which uses an 8:1:1 scaffold split with 3 random seeds.

Model	Test ROC-AUC	Protocol	Compute (downstream user)	Source
Logistic regression / ECFP	0.702 ± 0.018	MoleculeNet scaffold split	□ \$1	Wu et al. 2018, ESI Table S2
KernelSVM / ECFP	0.792	MoleculeNet scaffold split	~\$1	Wu et al. 2018, ESI Table S2
GraphConv	0.763 ± 0.016	MoleculeNet scaffold split	~\$1–10	Wu et al. 2018, ESI Table S2
AttentiveFP (Xiong et al., 2019)	0.757 ± 0.014	8:1:1 scaffold, 3 seeds	~\$10s	Re-run in Zhou et al. 2023, Table 1
D-MPNN / Chemprop (Yang et al., 2019)	0.771 ± 0.005	8:1:1 scaffold, 3 seeds	~\$10–100	Re-run in Zhou et al. 2023, Table 1

Model	Test ROC-AUC	Protocol	Compute (downstream user)	Source
GROVER_base (Rong et al., 2020)	0.625 $\pm$ 0.009	8:1:1 scaffold, 3 seeds	~15K V100-hr pretraining + finetune	Zhou et al. 2023, Table 1
GROVER_large (Rong et al., 2020)	0.682 $\pm$ 0.011	8:1:1 scaffold, 3 seeds	~24K V100-hr pretraining + finetune	Zhou et al. 2023, Table 1
This work (GNN + MolFormer ensemble)	0.806 $\pm$ 0.018	5-fold scaffold CV	\$0 (uses public MolFormer-XL)	this paper, §7.1
Uni-Mol (Zhou et al., 2023)	0.808 $\pm$ 0.003	8:1:1 scaffold, 3 seeds	~160 V100-hr pretraining (□ \$1K)	Zhou et al. 2023, Table 1

Two protocol caveats are essential. First, the Uni-Mol/GEM reproductions in Table 1 above use an 8:1:1 scaffold split with 3 random seeds; my numbers come from 5-fold scaffold CV. These protocols are not strictly comparable — but each test set is genuinely scaffold-held-out, so a 0.002–0.005 AUC difference is well within protocol-induced noise. Second, GROVER’s HIV scaffold numbers (0.625, 0.682) are from Uni-Mol’s reproduction, **not from the original GROVER paper, which does not report HIV in its main tables**. They appear here for completeness; readers should be aware that the GROVER authors did not themselves claim a strong HIV result.

Substantive observations:

1. The current SOTA on this benchmark (Uni-Mol, 0.808 $\pm$ 0.003) is statistically tied with my zero-cost ensemble (0.806 $\pm$ 0.018, with 95% fold-mean CI [0.791, 0.819]).
2. GROVER — by far the heaviest pretraining investment among 2D models — performs *worse* than a logistic-regression baseline on HIV scaffold-split. Pretraining scale is not a reliable predictor of downstream AUC on this benchmark.
3. MAT (Maziarka et al., 2020) is sometimes cited with an HIV result; the original paper does not benchmark HIV. I therefore omit it from this table.

8. Compute cost analysis

I document compute costs explicitly. The honest comparison is more nuanced than I initially framed it: Uni-Mol’s pretraining is *not* the multi-hundred-thousand-dollar cluster run that informal narratives sometimes suggest. The asymmetry is real but smaller, and the more striking observation is the matched downstream cost.

My pipeline (downstream cost): - GNN training: ~1 hour per fold on Apple Silicon MPS (consumer MacBook Pro), 5 folds = ~5 hours wall-clock, \$0. - MolFormer fine-tuning: ~1 hour per fold on Google Colab free-tier T4, 5 folds = ~5 GPU-hours of free-tier time, \$0. - Stacker fitting: ~10 minutes on consumer hardware, \$0. - Total: ~10 hours wall-clock distributed across free services. **Total downstream compute cost: \$0.**

Uni-Mol pretraining (per Zhou et al. 2023, Appendix C.1): - Molecular pretraining: 8x V100 GPUs x ~20 hours = ~160 V100-GPU-hours. - At AWS p3.2xlarge on-demand rates (\$3.06/V100-hour): \$490. - At spot/reserved rates: **\$200**. - The full Uni-Mol release also includes a pocket model (~544 V100-hours, another ~\$1,700 on-demand). For HIV property prediction only the molecular model is used. - Including their reported finetuning sweep (4x V100 GPUs x several days), realistic total replication cost is **low single-digit thousands of dollars**.

Cost ratio (corrected). My pipeline’s marginal downstream cost is \$0; Uni-Mol’s pretraining cost is roughly **\$500–\$2,000**. The pretrained MolFormer-XL checkpoint that I fine-tune from is itself the result of substantial upstream compute (1.1B-molecule pretraining at IBM); my claim is not “machine learning is free” but rather “the marginal cost to a downstream user has collapsed to zero”.

The cleaner observation is **performance**, not cost: my zero-marginal-cost MolFormer-XL fine-tuning *matches* Uni-Mol’s HIV result (per-fold mean 0.806 ± 0.018 vs 0.808 ± 0.003) without using any 3D conformers, custom pretraining, or institutional GPU access; the GNN+MolFormer stacker further yields a small but significant additive gain (OOF AUC 0.865 vs MolFormer-alone OOF AUC 0.856, paired bootstrap $p = 0.0015$). The corresponding cost-per-AUC-point claim from earlier drafts ($10^5\times$ advantage) was overstated by roughly two orders of magnitude and is removed.

9. Limitations

I bound my claims explicitly.

1. **Scaffold-split, retrospective evaluation only.** I performed no prospective wet-lab validation. Real drug discovery requires assay-confirmed hits, not held-out AUC.
2. **Single benchmark.** I report all results on MoleculeNet HIV. Generalization to other MoleculeNet tasks (BBBP, BACE, ClinTox, etc.) is plausible but not yet demonstrated. Multi-task evaluation is the obvious next step.
3. **Public pretraining is upstream cost.** While *my* pipeline incurs \$0 marginal cost, MolFormer-XL’s pretraining was itself expensive. My claim is that this cost has been *amortized* by IBM’s public release; the marginal cost to a downstream user is now zero. A fully bootstrapped \$0 result starting from random weights remains out of reach.
4. **No new architecture.** Both base models are public; my contribution is methodological (ensembling, scaffold-honest evaluation, threshold calibration) rather than architectural.
5. **Single-benchmark tie, not generalized parity.** The 0.806 vs 0.808 result places my ensemble within statistical noise of Uni-Mol *on HIV*. Uni-Mol reports stronger absolute numbers than my work on several other MoleculeNet datasets (e.g. BBBP, BACE, ClinTox); my tie is not a claim that 2D ensembling matches 3D pretraining in general. HIV may be unusually well-suited to the 2D, sequence-based features that dominate my ensemble.
6. **Protocol mismatch with primary baselines.** The Uni-Mol/GEM Table-1 reproductions I compare against use 8:1:1 scaffold split with 3 random seeds; I use 5-fold scaffold CV. These protocols are similar in spirit but not identical. A fully apples-to-apples comparison would require re-running my model under the 8:1:1/3-seed protocol or re-running Uni-Mol under 5-fold CV; both are deferred future work.

7. **Earlier-draft errors disclosed.** Earlier drafts of this manuscript reported a Uni-Mol pre-training cost of “\$100,000+” and a Uni-Mol HIV AUC of 0.831. Both were wrong: the cost is roughly \$500–\$2,000 (Zhou et al. 2023, Appendix C.1) and the HIV AUC is 0.808 (Table 1). The 0.831 figure was a transcription error from GROVER’s BACE result. Earlier drafts also cited an HIV AUC for MAT (Maziarka et al. 2020), which never benchmarked HIV. The headline result (0.806 ± 0.018) is unaffected by these corrections; the cost-asymmetry argument was substantially weakened, which I have reflected throughout.
-

10. Conclusion

This paper documents that, as of 2026, a careful ensemble of two publicly-available molecular models — combined through honest out-of-fold stacking, evaluated on scaffold-split CV, and calibrated through principled threshold tuning — recovers competitive performance on MoleculeNet HIV at zero dollars of downstream compute cost. The fine-tuned MolFormer-XL component achieves a per-fold test ROC-AUC of **0.806 ± 0.018** , statistically tied with Uni-Mol’s 0.808 ± 0.003 (the current SOTA on this benchmark); the GNN + MolFormer logistic stacker yields an out-of-fold ensemble AUC of **0.865**, a small but statistically significant additive gain (paired bootstrap $p = 0.0015$ vs. MolFormer-alone). The headline finding is the *performance match* without 3D conformers or institutional pretraining, not a several-orders-of-magnitude cost asymmetry.

I interpret this as evidence that, for binary molecular property prediction tasks where strong public pretrained checkpoints exist, the marginal value of bespoke pretraining has narrowed substantially on at least one standard benchmark. Engineering effort applied to ensemble design and evaluation methodology now recovers competitive AUC at zero downstream compute cost. I do not claim this generalizes to all MoleculeNet tasks; multi-benchmark replication is the natural next step.

The methodological implication is that low-resource molecular ML — accessible to students, independent researchers, and academic groups without GPU clusters — is now a credible path to competitive performance on standard benchmarks. I release all code, splits, and trained checkpoints publicly to support reproduction.

11. Reproducibility

I release all code, configuration, and splits at:

Source code: <https://github.com/v659/HIV-drug-discovery> (MIT licensed)

Trained model checkpoints, per-fold normalization statistics, the fitted ensemble stacker, and the preprocessed graph cache are archived on Zenodo with a citable DOI:

Pretrained artifacts (Zenodo): <https://doi.org/10.5281/zenodo.19946459>

The full pipeline (GNN training + MolFormer fine-tuning + stacker fitting + inference) reproduces from a clean clone in approximately 10 hours of wall-clock time on Google Colab free-tier and

Apple Silicon, at \$0 cost. To reproduce inference only without retraining, download the Zenodo bundle and unzip into `src/`.

I provide a complete reproduction recipe, including the exact pip environment and version-pinned dependencies (notably `transformers==4.46.3` to maintain MolFormer-XL compatibility), in `README.md`.

Acknowledgments

This work used Google Colaboratory’s free GPU tier and IBM’s publicly-released MolFormer-XL pretrained checkpoint. I acknowledge that the existence of these public resources is what makes a \$0 compute budget feasible. I worked on this project as an independent researcher, without institutional GPU access or compute budget.

AI assistance disclosure. I used Anthropic’s Claude (Claude Code) for manuscript drafting, code review, and figure-generation tooling. All experimental design, training and evaluation runs, statistical analyses, and final claims are mine; the trained models and reported numbers are produced by the code in this repository, not generated by an AI.

References

1. Bemis, G. W. & Murcko, M. A. (1996). The Properties of Known Drugs. 1. Molecular Frameworks. *Journal of Medicinal Chemistry*, 39(15), 2887–2893. [doi:10.1021/jm9602928](https://doi.org/10.1021/jm9602928)
2. Brody, S., Alon, U., & Yahav, E. (2022). How Attentive are Graph Attention Networks? *International Conference on Learning Representations (ICLR)*. [arXiv:2105.14491](https://arxiv.org/abs/2105.14491)
3. Chithrananda, S., Grand, G., & Ramsundar, B. (2020). ChemBERTa: Large-Scale Self-Supervised Pretraining for Molecular Property Prediction. *Machine Learning for Molecules Workshop, NeurIPS 2020*. [arXiv:2010.09885](https://arxiv.org/abs/2010.09885)
4. Duvenaud, D., Maclaurin, D., Aguilera-Iparraguirre, J., Gómez-Bombarelli, R., Hirzel, T., Aspuru-Guzik, A., & Adams, R. P. (2015). Convolutional Networks on Graphs for Learning Molecular Fingerprints. *Advances in Neural Information Processing Systems (NeurIPS)*, 28. [arXiv:1509.09292](https://arxiv.org/abs/1509.09292)
5. Kearnes, S., McCloskey, K., Berndl, M., Pande, V., & Riley, P. (2016). Molecular Graph Convolutions: Moving Beyond Fingerprints. *Journal of Computer-Aided Molecular Design*, 30(8), 595–608. [doi:10.1007/s10822-016-9938-8](https://doi.org/10.1007/s10822-016-9938-8)
6. Landrum, G. (2024). RDKit: Open-source cheminformatics. <https://www.rdkit.org>. (Used for Murcko scaffold extraction, ECFP4 fingerprints, and global descriptor computation.)
7. Lin, T.-Y., Goyal, P., Girshick, R., He, K., & Dollár, P. (2017). Focal Loss for Dense Object Detection. *IEEE International Conference on Computer Vision (ICCV)*. [arXiv:1708.02002](https://arxiv.org/abs/1708.02002)




Figure 1: **Figure 1.** Per-fold scaffold-held-out test AUC for the GNN and MolFormer-XL across all five folds, with mean reference lines.




Figure 2: **Figure 2.** ROC curves on n = 24,391 out-of-fold predictions for Tanimoto-NN, GNN, MolFormer-XL, and the final stacker.




Figure 3: **Figure 3.** Precision-Recall curves on the same OOF predictions, with the 3.7%-base-rate random reference line.




Figure 4: **Figure 4.** Reliability diagram (quantile-binned) comparing raw MolFormer-XL probabilities against the calibrated stacker output. Axes are zoomed to $[0, 0.5] \times [0, 0.3]$ because at a 3.7% base rate the upper-right of the unit square is unpopulated.




Figure 5: **Figure 5.** Recall, precision, and F1 vs decision threshold on the OOF stacker output, with Youden's J, base-rate, and F1-max thresholds marked.

Appendix A — Threshold sweep on held-out mixed set

Threshold	TP	FP	TN	FN	Recall	Precision	F1	FPR
0.020	98	309	691	2	0.980	0.241	0.387	0.309
0.043	90	125	875	10	0.900	0.419	0.571	0.125
(Youden)								
0.100	76	48	952	24	0.760	0.613	0.679	0.048
0.173	68	21	979	32	0.680	0.764	0.720	0.021
(base-rate)								
0.220	66	14	986	34	0.660	0.825	0.733	0.014
(F1-max)								
0.300	60	9	991	40	0.600	0.870	0.710	0.009
0.500	49	5	995	51	0.490	0.907	0.636	0.005

Appendix B — Cost itemization

Component	Hardware	Wall time	Cost
GNN training, fold 0	Apple M-series MPS	~1 hr	\$0
GNN training, folds 1–4	Apple M-series MPS	~4 hr	\$0
MolFormer fine-tuning, fold 0	Colab T4 (free)	~1 hr	\$0
MolFormer fine-tuning, folds 1–4	Colab T4 (free)	~4 hr	\$0
Stacker fitting	Apple M-series MPS	~10 min	\$0
Inference + evaluation	Apple M-series MPS	~5 min	\$0
Total		~10 hr	\$0